

Vivian Nguyen

San Jose, CA • [linkedin.com/in/vivian-k-nguyen/](https://www.linkedin.com/in/vivian-k-nguyen/)

EDUCATION

University of California, Santa Cruz

Santa Cruz, CA

Bachelor of Science (B.S) in Computer Science

Bachelor of Science (B.S) in Technology and Information Management

SKILLS

Languages & Frameworks: Python, SQL, C/C++, JavaScript, Pandas, FastAPI, Streamlit, XGBoost, Anthropic API

Tools & Practices: GitHub, PostgreSQL, Vercel, Prompt Engineering, Model Evaluation, Claude Code, Agent Workflows

WORK EXPERIENCE

Kool Quarters

San Jose, CA

AI Engineering Intern

Jul 2026 - Present

- Construct automated data collection infrastructure and deploy FastAPI services (Docker, Railway, Vercel, Supabase) that feed real estate market analysis platforms, processing high volumes of unstructured web data into clean formats
- Standardize ingestion workflows across external sources including Census, county assessor, and MLS databases, guaranteeing strict data fidelity and per-value provenance for downstream valuation models and evaluation algorithms
- Optimize extraction system reliability with fallback protocols, rate limits, error monitoring, and 450+ automated tests to prevent silent failures across active data pipelines

IDXExchange

San Jose, CA

Data Science Intern

May 2026 - Sep 2026

- Built an end-to-end ML pipeline in Python that ingests CRMLS sold-transaction data via FTP, validates schemas, and handles missing values and categorical encoding, producing reproducible, production-ready datasets
- Implemented spatial feature engineering with GeoPandas, joining property coordinates against California school-district boundaries to enrich location-sensitive valuation models
- Structured the project as modular preprocessing, training, evaluation, and prediction scripts with a documented README, and deployed a Streamlit app that serves XGBoost predictions from user inputs

Misfit Lab

Santa Cruz, CA

Research Assistant

Jun 2025 - Jan 2026

- Built a syllabus-driven Slack productivity bot in Python that integrates through the Slack API, parses course schedules into a SQL database for persistent storage, and dispatches scheduled deadline reminders and task-splitting suggestions to student project channels
- Tracked and documented user behaviors and bot responses across 50+ test scenarios, building structured data pipelines to measure what drove engagement, click-through, and follow-through on deadlines
- Facilitated remote focus-group feedback sessions and managed end-to-end participant recruitment to turn qualitative findings into feature changes

PROJECTS

Z-Scraper | Python, FastAPI, REST APIs, Supabase, SQLite, Zillow & CRMLS Data

A Python service that extracts and normalizes unstructured property listings into structured, traceable market-evaluation records for real estate investors.

- Develop scalable Python listing extraction services that parse, clean, and normalize unstructured property URLs into unified schemas for real estate market evaluation
- Architect audit-ready data schemas with field-level provenance tracking (source, vintage, confidence) and explicit NULL handling to maintain data integrity for downstream investment models
- Implement automated resilience protocols including retries, rate limiting, and failure logging to ensure property records reliably join with Census and county assessor datasets without silent data loss

LinkedIn Blogger | Python, Anthropic Claude API, Flask, LinkedIn API, GitHub API

An LLM pipeline converting a week of work notes and GitHub activity into LinkedIn posts, gated by human review.

- Created a multi-step LLM pipeline on the Anthropic SDK: idea generation, skeleton drafting with author-voice gaps, and a structured-JSON review pass merged with deterministic checks so length and factual flags are never left to the model
- Built a shared Anthropic API client that retries on rate-limit and overload errors, parses JSON responses with clear error messages, and generates image alt text with Claude vision
- Designed the approval gate and editable publish queue so nothing posts without a human decision, with version history and word-level diffs across every draft revision